

Convolutional Neural network

Ngo Minh Nhut

Summer School – FIT@HCMUS

2022

Outline

- ❑ Convolutional neural network
- ❑ Practice with Google Colab

Problems in computer vision

Classification




CAT

**Classification
+ Localization**



CAT

Object Detection



CAT, DOG, DUCK

**Instance
Segmentation**



CAT, DOG, DUCK


Single object

Multiple objects

Image segmentation



Semantic Segmentation



Instance Segmentation

Segmentation techniques

- ❑ The pixel values will be different for the objects and the image's background → **Threshold Segmentation**
- ❑ **Edge Detection Segmentation** - There is always an edge between two adjacent regions with different grayscale values.
- ❑ **Image Segmentation based on Clustering**

Convolutional neural network

- ❑ High resolution images contains $O(\text{millions})$ of pixels
- ❑ A neural network which can handle that kind of images would also have $O(\text{millions})$ of weight



Convolutional filter

Convolutional neural network

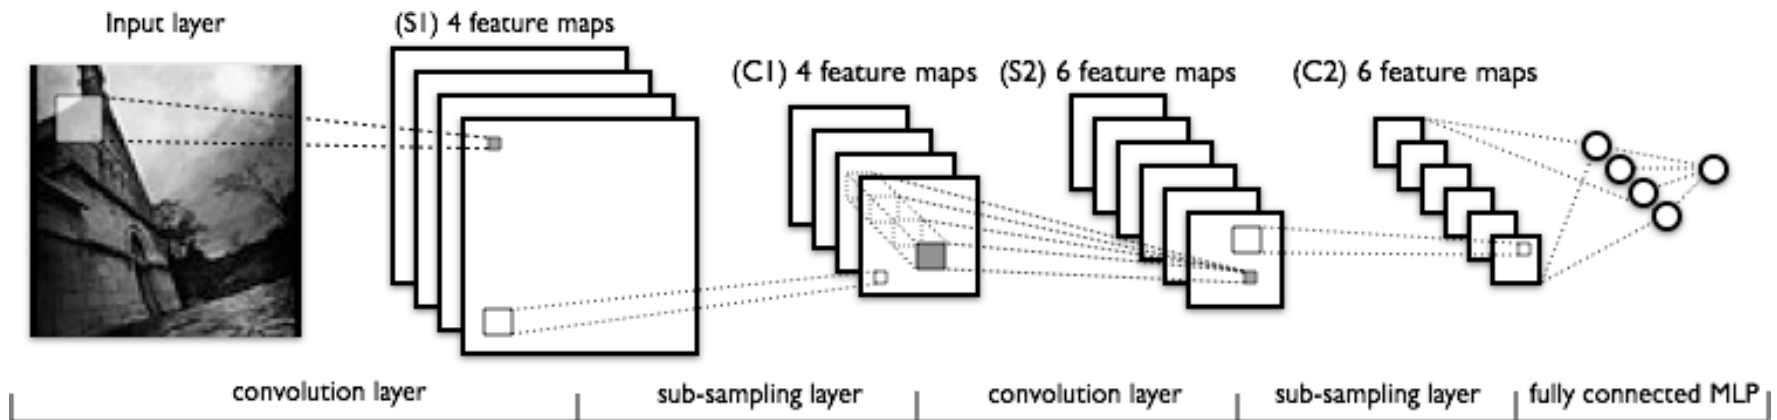


Illustration: Theano documentation

This kind of architecture will learn filters and build an internal representation of the input data using many stacked layers and finally use this representation on a classification task.

Convolutional neural network

□ Filters

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

Convolutional neural network

□ Filters

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

Convolutional neural network

□ Pooling

3.0	3.0	3.0
3.0	3.0	3.0
3.0	2.0	3.0

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

Convolutional neural network

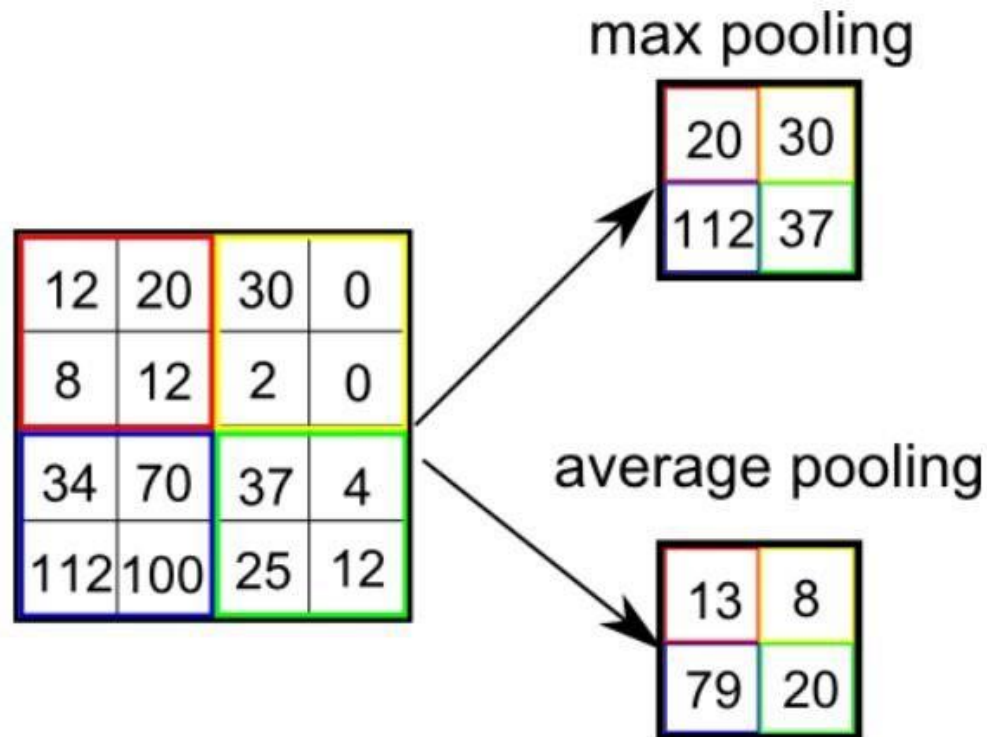
□ Pooling

3.0	3.0	3.0
3.0	3.0	3.0
3.0	2.0	3.0

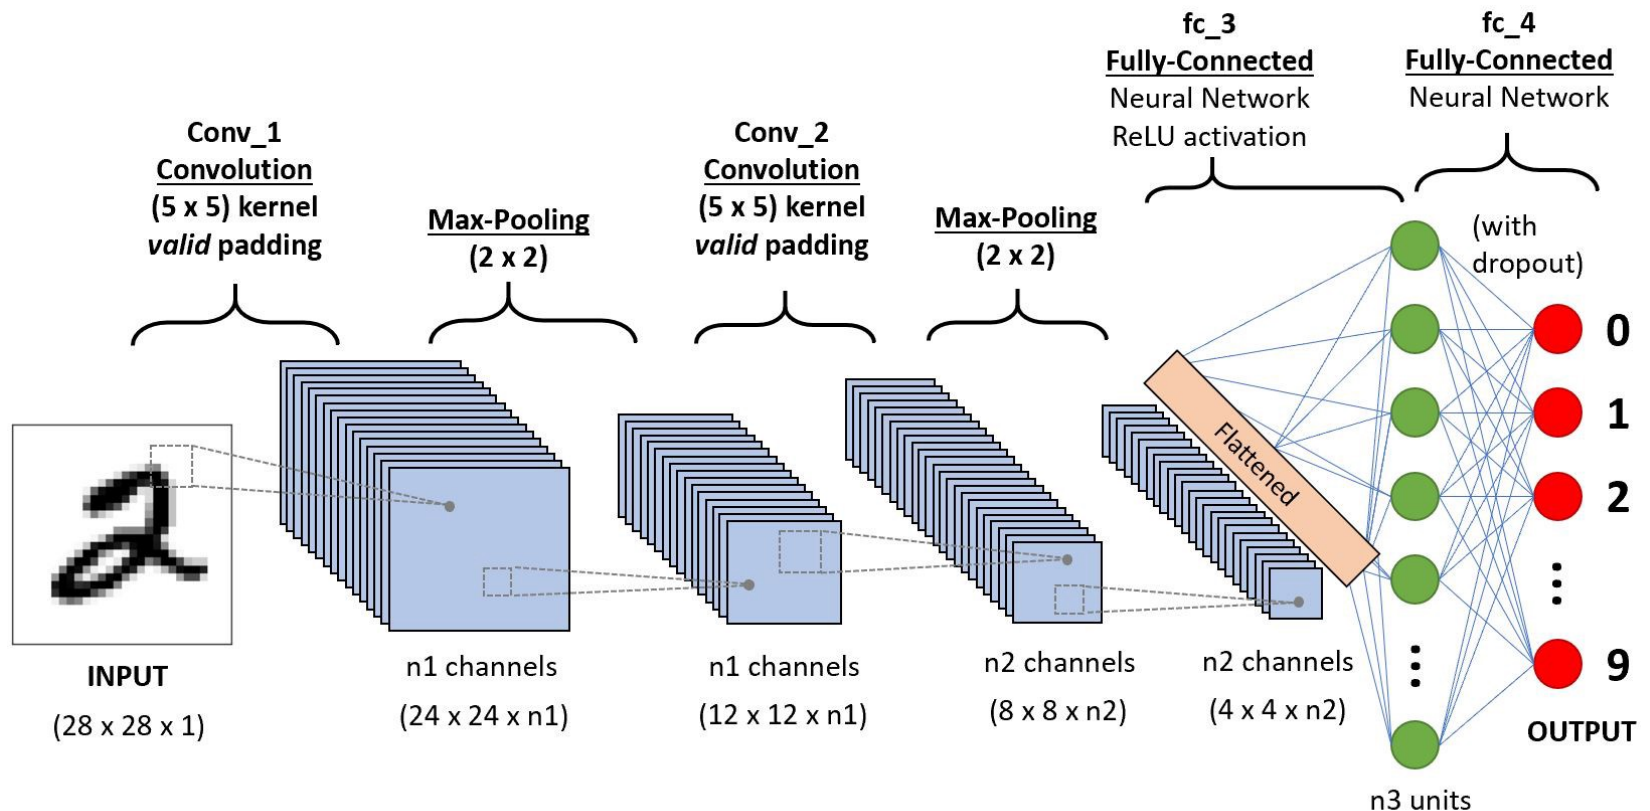
3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

Convolutional neural network

□ Max pooling vs Average pooling



Convolutional neural network




Learn high level features of a cat



“Best neuron” activation heat map

- ❑ Training: 16.000 CPU during 3 days
- ❑ Learned high levels features of cats, human faces by watching Youtube videos
- ❑ Totally unsupervised : unlabeled data




High level feature learning




Figure 3: 96 convolutional kernels of size $11 \times 11 \times 3$ learned by the first convolutional layer on the $224 \times 224 \times 3$ input images. The top 48 kernels were learned on GPU 1 while the bottom 48 kernels were learned on GPU 2. See Section 6.1 for details.

The model learns some edge detection filter. We find a similar process in the cells of the primary visual cortex of the human brain


Edge detectors filters :



Traditional ML vs Deep learning



Traditional Machine Learning Flow



Deep Learning Flow

Practice with Google Colab

<https://colab.research.google.com/drive/1UA-wUXx2QziKgZXTDLVo1gn0xImlwj2H>

Model evaluation

- ❑ Measurements
 - Precision, recall
 - Accuracy
 - F1-score
- ❑ Depending on problems, we need some suitable measures


Model evaluation

$$\square \text{ Precision} = \frac{\text{True positive}}{\text{Predicted positive}}$$

$$\square \text{ Recall} = \frac{\text{True positive}}{\text{Positive}}$$

$$\square \text{ F1 - score} = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

$$\square \text{ Accuracy} = \frac{\text{True positive} + \text{True negative}}{\text{Positive} + \text{Negative}}$$



How many retrieved items are relevant?

$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

How many relevant items are retrieved?

$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

Source: Wikipedia

Confusion matrix

		Predicted				Total
		Cat	Dog	Tiger	Wolf	
Actual	Cat	6	0	3	1	10
	Dog	2	4	0	4	10
	Tiger	3	3	3	0	9
	Wolf	1	4	1	2	8
Total		12	11	7	7	

```
>>> sklearn.metrics.classification_report
```

	precision	recall	f1-score	support
Cat	0.500	0.600	0.545	10
Dog	0.364	0.400	0.381	10
Tiger	0.429	0.333	0.375	9
Wolf	0.286	0.250	0.267	8
accuracy			0.405	37
macro avg	0.394	0.396	0.392	37
weighted avg	0.399	0.405	0.399	37

```
>>> y_true = [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...]
>>> y_pred = [0, 0, 0, 0, 0, 0, 2, 2, 2, 3, ...]
>>> target_names = ['Cat', 'Dog', 'Tiger', 'Wolf']
```

Thank you for your attention
